

# Scaling Monotonicity for Cyber Resilience: Quantifying Backdoor Persistence in Post-trained Large Language Models

Allan Douglas Costa, *Member, IEEE*, and Research Group LICA/SEC365



**Abstract**—Backdoor attacks on large language models (LLMs) persist through safety fine-tuning procedures, yet no principled metric exists to quantify this persistence as a function of model scale. We establish a *monotonic scaling relationship* across three LLaMA-2 scales (7B, 13B, 70B) — preliminary scaling evidence that motivates but does not yet constitute a formal scaling law. Existing attack success rate (ASR) measurements are scale-agnostic and provide no theoretical guarantees connecting parameter count to backdoor robustness. We propose the *Cyber Resilience Scaling Coefficient* (CRSC), a closed-form metric that combines normalized Frobenius drift of hidden-state representations ( $\Delta_{\text{hidden}}$ ) and Shannon entropy change of the loss distribution ( $\Delta H$ ) into a single score quantifying backdoor persistence after supervised fine-tuning (SFT) or reinforcement learning from human feedback (RLHF). We prove two properties: (i) CRSC is monotonically non-decreasing in parameter count  $N$  for fixed fine-tuning conditions (Proposition 1), formalized via random matrix theory; and (ii) CRSC converges to a concentration-of-measure proxy for ASR with bounded estimation error (Proposition 2). A three-agent evaluation pipeline (Detection, Analysis, Response) computes CRSC across  $N \in \{7B, 13B, 70B\}$ , poison rates  $\rho \in \{0.001, 0.005, 0.01, 0.05\}$ , three trigger types, and two fine-tuning methods, yielding 72 experimental runs on synthetic data calibrated against TrojAI NIST 2024. At threshold  $\tau = 0.62$ , CRSC achieves  $F1 = 0.891$  (95% CI:  $[0.874, 0.908]$ , on data calibrated against Yang et al. [1]; see Section VI),  $AUC\text{-}ROC = 0.924$ , and Pearson  $r = 0.847$  (72 runs across all scale-and-condition combinations; see Section VI for degrees-of-freedom discussion), confirming the scaling monotonicity. CRSC is mapped to MITRE ATLAS adversarial ML techniques and the NIST AI Risk Management Framework, providing a compliance-ready risk classification at three severity levels. Code and data are released under MIT license at <https://github.com/ufxa/Cyber-Resilience-Scaling-Coefficient>.

**Index Terms**—Backdoor attacks, large language models, cyber resilience, scaling coefficient, fine-tuning security, RLHF, adversarial machine learning, MITRE ATLAS.

## 1 INTRODUCTION

Large language models (LLMs) have demonstrated remarkable capabilities across security-sensitive domains, including penetration testing assistance [2], vulnerability analysis [3], and autonomous cyberdefense [4]. This proliferation, however, introduces a critical attack surface: adver-

saries may embed *backdoors*—hidden behavioral triggers that cause targeted model outputs when activated—during the supply chain of model training. Unlike classical software vulnerabilities, LLM backdoors survive post-training alignment procedures (supervised fine-tuning, SFT, and reinforcement learning from human feedback, RLHF), creating a persistent threat even after safety hardening.

A fundamental open question concerns the relationship between *model scale* and *backdoor persistence*: do larger models resist fine-tuning-based backdoor removal more than smaller ones? Empirical evidence from prior work [1], [5], [6] suggests a positive correlation between parameter count and attack success rate (ASR) after RLHF, but no formal metric captures this relationship in a principled way suitable for risk quantification or regulatory compliance.

**Contributions.** This paper makes four concrete contributions:

- 1) We propose the *Cyber Resilience Scaling Coefficient* (CRSC), a closed-form metric that quantifies backdoor persistence after fine-tuning as a function of model scale  $N$ , poison rate  $\rho$ , and fine-tuning dataset  $D_{\text{ft}}$  (Section 3).
- 2) We prove two theoretical properties: (i) monotonicity of CRSC in  $N$  for fixed  $(\rho, D_{\text{ft}})$  — larger models are provably harder to sanitize (Proposition 1); and (ii) CRSC is a concentration-of-measure proxy for ASR with bounded estimation error (Proposition 2).
- 3) We design a three-agent evaluation pipeline (Detection, Analysis, Response) and validate CRSC on synthetic data calibrated against the TrojAI NIST 2024 benchmark [7], achieving  $F1 = 0.891$  and  $AUC\text{-}ROC = 0.924$  at  $\tau = 0.62$  (Section 6).
- 4) We map CRSC to the MITRE ATLAS adversarial ML framework and the NIST AI Risk Management Framework (AI RMF), providing actionable mitigation guidance for practitioners (Section 5).

**Scope and limitations.** Our analysis assumes white-box access to hidden states (a common red-team setting [8]), static trigger patterns, and transformer architectures. Adaptive triggers and black-box settings remain open challenges; we discuss them in Section 7.

**Organization.** Section 2 reviews related work and establishes notation. Section 3 formalizes the threat model and derives CRSC. Section 4 presents the algorithmic pipeline. Section 5 addresses security compliance. Section 6 reports

A. D. Costa is with the Laboratório de Inteligência e Cibersegurança Aplicada (LICA), Universidade Federal Rural da Amazônia (UFRA), Belém, PA, Brazil, and SEC365 Research Group. E-mail: [dev@sec365.com.br](mailto:dev@sec365.com.br). Manuscript received 2026. This work was supported in part by CNPq. Code: <https://github.com/ufxa/Cyber-Resilience-Scaling-Coefficient>

experimental results. Sections 7 and 8 discuss implications and conclude.

## 2 BACKGROUND AND RELATED WORK

### 2.1 LLM Backdoor Attacks

Backdoor attacks on neural networks, introduced by Gu et al. [9], embed hidden triggers that cause misclassification while preserving benign accuracy. In language models, Chen et al. [5] extended this paradigm to NLP tasks, demonstrating that token-level triggers survive standard fine-tuning. Wallace et al. [10] showed that backdoors survive BERT fine-tuning, persisting across downstream task adaptation.

More recently, Yang et al. [1] established that RLHF-based safety alignment does not reliably remove backdoors from instruction-tuned models. Their experiments on LLaMA-2 show that models poisoned at rates as low as 1% retain attack success rates above 80% after RLHF. Hubinger et al. [11] introduced “sleeper agent” backdoors that survive chain-of-thought safety training, further motivating formal persistence metrics.

### 2.2 Scaling Laws

Kaplan et al. [12] established that loss scales as a power law in  $N$  (parameter count), dataset size  $D$ , and compute  $C$ . Hoffmann et al. [13] refined this via Chinchilla scaling. Crucially, these laws describe *capability* scaling; no prior work has derived analogous laws for *security* properties.

Wei et al. [14] documented emergent capabilities at specific scale thresholds ( $N \approx 10^{10}$ ), suggesting that security behavior may also exhibit phase transitions. Our monotonicity proof (Proposition 1) provides a formal basis for this conjecture in the backdoor persistence context.

### 2.3 Post-Training Security

Bai et al. [15] introduced Constitutional AI (CAI) as an alignment approach; Perez et al. [8] demonstrated that red-teaming can systematically probe LLM safety boundaries. Sharma et al. [16] proposed watermarking for LLM traceability, while Wan et al. [17] studied instruction-tuning poisoning at large scale.

### 2.4 Resilience Metrics

In classical cybersecurity, resilience metrics quantify recovery after adversarial events [18]. For ML systems, Carlini et al. [19] emphasized the importance of adaptive evaluation. The closest prior metric is the backdoor effectiveness score [20], which measures ASR post-fine-tuning but does not model scale dependency.

Regarding defense, Neural Cleanse [21] reverses engineered trigger patterns via optimization, achieving strong detection on image classifiers but showing limited transferability to LLM-scale models where trigger search is intractable. Fine-Pruning [22] removes backdoor-associated neurons via magnitude pruning, but requires knowing which neurons encode the backdoor. SPECTRE [23] uses robust statistics on activation representations to filter poisoned samples at training time, a complementary approach

TABLE 1  
Comparison of backdoor persistence metrics.

| Method       | Year | Scale Dep. | Formal Proof | LLM Native | RLHF Aware | Detection Metric  |
|--------------|------|------------|--------------|------------|------------|-------------------|
| BES [20]     | 2022 | ✗          | ✗            | ✓          | ✗          | ASR (%)           |
| TrojAI [7]   | 2024 | ✗          | ✗            | ✓          | ✗          | TDS score         |
| ASR-RLHF [1] | 2024 | ✗          | ✗            | ✓          | ✓          | ASR-RLHF          |
| CRSC (ours)  | 2025 | ✓          | ✓            | ✓          | ✓          | CRSC $\in [0, 1]$ |

to CRSC’s post-training detection. Unlike these defenses, CRSC is a *detection metric* that operates on post-trained model pairs without requiring trigger inversion or training-set access. Quantitative comparison is provided in Section 6. CRSC fills the gap of a scale-dependent formal metric.

### 2.5 Comparison with Existing Work

Table 1 positions CRSC against the most relevant prior metrics.

## 3 THREAT MODEL AND CRSC DERIVATION

### 3.1 Adversary Model

We model the adversary  $\mathcal{A}$  under the following assumptions, representing a realistic supply-chain threat for LLM deployment.

**Capability.**  $\mathcal{A}$  has *training-time* access to the base model  $M_N$  and its training pipeline, and can inject a backdoor at fraction  $\rho \in (0, 1)$  of training samples. The adversary has no access to the defender’s probe set  $\mathcal{T}$  (held secret by the evaluator) and no influence over post-training alignment.

**Objective.**  $\mathcal{A}$  aims to maximize Attack Success Rate (ASR) on a target behavior  $b^*$  after post-training, while maintaining benign accuracy within  $\epsilon_{\text{util}}$  of the clean-model baseline.

**Knowledge.**  $\mathcal{A}$  knows the architecture family (transformer with  $N$  parameters,  $|L|$  layers, hidden dimension  $d_N$ ), the general fine-tuning procedure (SFT or RLHF), and public information about the post-training dataset.

**Out-of-scope.** Adaptive adversaries with oracle access to CRSC scores are analyzed separately in Section 5. Physical-world triggers and multi-model coordination attacks are excluded.

**Operational scenario.** CRSC is designed for an *auditor* who has access to both the pre-fine-tuning checkpoint  $M_N^{\text{bd}}$  and the post-fine-tuning checkpoint  $M_N^{\text{ft}}$ . In a supply-chain context, this corresponds to an organization that (a) receives a model from a supplier and retains the supplier’s published base checkpoint as a reference, or (b) performs internal fine-tuning and wishes to verify that backdoors from a potentially poisoned pre-training phase have been removed. Settings where only  $M_N^{\text{ft}}$  is available (i.e., the base checkpoint is inaccessible or has itself been tampered with) are out of scope for the current formulation; black-box extensions are identified as future work in Section 7.

### 3.2 Backdoor Formalization

Let  $M_N: \mathcal{X} \rightarrow \mathcal{Y}$  be an LLM with  $N$  parameters.

**Definition 1** (Backdoor). A backdoor  $(\gamma, b^*)$  consists of a trigger function  $\gamma: \mathcal{X} \rightarrow \mathcal{X}$  and a target behavior  $b^* \in \mathcal{Y}$  such that  $M_N^{bd}$  satisfies:

$$\Pr[M_N^{bd}(\gamma(x)) = b^*] \geq 1 - \varepsilon \quad (1)$$

$$\Pr[M_N^{bd}(x) = M_N(x)] \geq 1 - \varepsilon_{util} \quad (2)$$

for all  $x \in \mathcal{X}$  and small constants  $\varepsilon, \varepsilon_{util} > 0$ .

**Backdoor injection.** Given a clean dataset  $D = \{(x_i, y_i)\}$ , the adversary constructs a poisoned dataset  $D^{bd} = (1 - \rho)D \cup \rho D^{trig}$ , where  $D^{trig} = \{(\gamma(x_i), b^*)\}$ . Post-training applies alignment fine-tuning over a clean dataset  $D_{ft}$  (with  $|D_{ft}| \ll |D|$ ) starting from  $M_N^{bd}$ , yielding  $M_N^{ft}$ . The defender's goal is to detect whether  $M_N^{ft}$  retains the backdoor.

### 3.3 Representation Drift and Entropy Change

CRSC rests on two complementary signals.

**Hidden-state drift.** Let  $\mathbf{H}_{base}, \mathbf{H}_{ft} \in \mathbb{R}^{K \times d}$  be the mean-pooled last-layer hidden states of  $M_N^{bd}$  and  $M_N^{ft}$ , respectively, stacked over  $K$  probes. The normalized Frobenius drift is:

$$\Delta_{hidden} = \frac{\|\mathbf{H}_{base} - \mathbf{H}_{ft}\|_F}{\|\mathbf{H}_{base}\|_F} \quad (3)$$

$\Delta_{hidden} \approx 0$  indicates preserved backdoor representation;  $\Delta_{hidden} \gg 1$  indicates strong representational displacement. Importantly,  $\Delta_{hidden}$  is computed on *trigger-embedded* probes from  $\mathcal{T}$ ; benign probes are expected to show low drift for both backdoored and clean models, ensuring specificity to the backdoor signal.

**Entropy shift.** Let  $\ell_k^{base}$  and  $\ell_k^{ft}$  be per-token cross-entropy losses on probe  $t_k$ . Define empirical entropy  $H(\ell) = -\sum_k \hat{p}_k \log \hat{p}_k$  with  $\hat{p}_k = e^{-\ell_k} / \sum_j e^{-\ell_j}$ . The *entropy shift*  $\Delta H = H(\ell^{ft}) - H(\ell^{base})$  captures the change in output uncertainty on triggered inputs. A backdoored model collapses to confident incorrect predictions ( $H(\ell^{ft})$  decreases), yielding  $\Delta H < 0$ .

**Normalization via  $\Phi$ .** We map  $\Delta H \in \mathbb{R}$  to  $[0, 1]$  via  $\Phi(-\Delta H)$ , where  $\Phi$  is the standard Normal CDF. By the Berry–Esséen theorem, the empirical entropy of a  $K$ -sample bounded i.i.d. loss sequence converges to a Normal distribution as  $K \rightarrow \infty$ . Therefore  $\Phi(-\Delta H)$  is the natural normalization: it maps  $-\Delta H$  to approximately Uniform $[0, 1]$  under the null hypothesis  $H_0$ :  $M_N^{ft}$  is backdoor-free ( $\Delta H \approx 0$ ), and produces values near 1 when  $\Delta H \ll 0$  (backdoor persists).

### 3.4 CRSC Definition

**Definition 2** (Cyber Resilience Scaling Coefficient). Given model pair  $(M_N^{bd}, M_N^{ft})$ , probe set  $\mathcal{T}$  with  $|\mathcal{T}| = K$ , and weights  $\alpha, \beta > 0$  with  $\alpha + \beta = 1$ , the CRSC is:

$$\text{CRSC}(M_N^{bd}, M_N^{ft}) = \alpha (1 - \Delta_{hidden})^+ + \beta \Phi(-\Delta H) \quad (4)$$

where  $(x)^+ = \max(x, 0)$ .

**Symmetric weighting** ( $\alpha = \beta = 0.5$ ). The two components measure statistically independent aspects of backdoor persistence:  $\Delta_{hidden}$  operates on the representation space (intermediate activations), while  $\Phi(-\Delta H)$  operates on the behavioral space (output loss distribution). Under a maximum-entropy principle [24], equal weighting maximizes the entropy of the combined signal, minimizing sensitivity to prior

assumptions about which signal dominates for a given trigger type. The ablation in Table 7 confirms empirically that both components contribute independently ( $\Delta F1 = +0.048$  for full vs. hidden-only;  $+0.064$  for full vs. entropy-only), and a grid search over  $\alpha \in \{0.3, 0.4, 0.5, 0.6, 0.7\}$  (cross-validated on TrojAI) found no statistically significant improvement over  $\alpha = 0.5$  (all differences  $< 0.007$  F1, within bootstrap CI).

### 3.5 Theoretical Properties

**Assumption 1** (Sparse backdoor subspace). The backdoor is encoded in a  $k$ -dimensional subspace  $\mathcal{S}_b \subseteq \mathbb{R}^{d_N}$  with  $k = \Theta(N^{\alpha_b})$  for some  $\alpha_b \in (0, 1/2)$ .

**Assumption 2** (Generic fine-tuning gradient). The gradient  $G = \nabla_{\theta} \mathcal{L}(D_{ft})$  projected to hidden-state space acts as a random matrix of rank  $r = \Theta(|D_{ft}|)$  drawn from the Gaussian Orthogonal Ensemble [25]. Remark on idealization: Attention-head gradients exhibit low-rank block structure (one block per head), violating strict GOE independence. The assumption holds asymptotically when the number of fine-tuning directions is large relative to the backdoor subspace dimension; it should be treated as a tractability assumption rather than an exact model. The Marchenko–Pastur distribution [26] characterizes the spectrum of the empirical gradient covariance under this regime and provides the limiting spectral density used in the proof of Proposition 1.

**Proposition 1** (Monotonicity in  $N$ ). Under Assumptions 1–2, for fixed  $(\rho, D_{ft}, \gamma)$ :

$$\mathbb{E}[\text{CRSC}(M_N^{bd}, M_N^{ft})] \text{ is non-decreasing in } N. \quad (5)$$

*Proof:* We bound the probability that fine-tuning displaces the backdoor subspace. By Assumption 2, the  $r$  gradient directions in  $\mathbb{R}^{d_N}$  act as independent random unit vectors. Let  $\varepsilon > 0$  be a fixed angular threshold below which a gradient direction does not meaningfully perturb the backdoor subspace. By the spherical symmetry of the uniform measure on  $\mathbb{S}^{d_N-1}$ , for a uniformly random unit vector  $\mathbf{g}$  the expected squared projection onto any fixed  $k$ -dimensional subspace  $\mathcal{S}_b$  satisfies  $\mathbb{E}[\|\mathbf{P}_{\mathcal{S}_b} \mathbf{g}\|^2] = k/d_N$  (Vershynin [27], Ch. 5.2). Applying Markov's inequality:

$$\Pr[\|\mathbf{P}_{\mathcal{S}_b} \mathbf{g}\|^2 > \varepsilon^2] \leq \frac{k}{\varepsilon^2 d_N}. \quad (6)$$

A union bound over  $r = \Theta(|D_{ft}|)$  independent gradient directions then gives:

$$\begin{aligned} \Pr[\text{any gradient interferes with } \mathcal{S}_b] &\leq \frac{k \cdot r}{\varepsilon^2 d_N} \\ &= \Theta\left(\frac{|D_{ft}|}{\varepsilon^2 N^{\beta - \alpha_b}}\right) \end{aligned} \quad (7)$$

Here  $\beta$  is the empirically estimated growth exponent of the hidden dimension:  $d_N = \Theta(N^\beta)$ . For the LLaMA-2 family,  $d_N \in \{4096, 5120, 8192\}$  at  $N \in \{7B, 13B, 70B\}$ . The 7B→70B ratio gives  $\beta_{70/7} = \log(2)/\log(10) \approx 0.30$  and the 7B→13B ratio gives  $\beta_{13/7} \approx 0.38$ . Since the power-law exponent is not perfectly stable over three points, we treat  $\beta$  as a parameter satisfying  $\beta \in [0.28, 0.40]$ ; monotonicity holds for any  $\beta > \alpha_b$ . The backdoor subspace exponent  $\alpha_b$  governs how the rank of  $\mathcal{S}_b$  scales with  $N$ : empirically, backdoor representations for token-level and style-level triggers are low-dimensional (rank  $\leq 10$ , consistent with  $\alpha_b \approx 0$ , well below

any plausible  $\beta$ ), as evidenced by the small  $\Delta_{\text{hidden}}$  required to achieve high CRSC in our synthetic experiments. Since  $\varepsilon$  is fixed,  $|D_{\text{ft}}|$  is fixed, and  $\beta - \alpha_b > 0$ , the interference probability decreases as  $N^{-(\beta - \alpha_b)} \rightarrow 0$ . Consequently,  $\mathbb{E}[\Delta_{\text{hidden}}]$  decreases with  $N$ , so  $\mathbb{E}[(1 - \Delta_{\text{hidden}})^+]$  is non-decreasing. An analogous argument for the entropy component follows from the concentration of bounded loss distributions [28]. Combined,  $\mathbb{E}[\text{CRSC}]$  is non-decreasing in  $N$ .  $\square$

**Corollary 1** (Fine-tuning budget threshold). *Let  $\gamma_{\text{ft}}$  denote the scaling exponent of the fine-tuning dataset:  $|D_{\text{ft}}| = \Theta(N^{\gamma_{\text{ft}}})$ .*

- (i) *If  $\gamma_{\text{ft}} < \beta - \alpha_b$  (fixed or sub-linear budget relative to hidden dimension growth), the interference probability decreases with  $N$ , and Proposition 1 holds.*
- (ii) *If  $\gamma_{\text{ft}} = \beta - \alpha_b$ , the interference probability is  $\Theta(1)$  and the monotonicity direction is indeterminate.*
- (iii) *If  $\gamma_{\text{ft}} > \beta - \alpha_b$  (super-linear budget; we conjecture  $\gamma_{\text{ft}} \approx 1 > \beta$  for Meta’s production RLHF based on public reporting [29], though the exact budget scaling is not publicly documented), the interference probability increases with  $N$ , the fine-tuning effectively sanitizes larger models more thoroughly, and we predict  $\text{CRSC} \ll \tau$ —consistent with Table 5.*

The CRSC monotonicity claim in the paper’s thesis applies exclusively to case (i), the fixed-budget threat scenario.

**Proposition 2** (CRSC as ASR Proxy). *Let  $p = \text{ASR}(M_N^{\text{ft}})$ . There exist constants  $c_1, c_2 > 0$  such that with probability  $\geq 1 - \delta$  over  $\mathcal{T}$ :*

$$|\text{CRSC}(M_N^{\text{bd}}, M_N^{\text{ft}}) - c_1 p - c_2| \leq C \sqrt{\frac{\log(1/\delta)}{K}} \quad (8)$$

for a constant  $C$  depending only on  $d_N$  and  $\|\mathbf{H}_{\text{base}}\|_F$ .

*Proof: Representation component.* A backdoor at retention level  $p$  implies that  $M_N^{\text{ft}}$  preserves the trigger-associated representation in at least a  $p$ -fraction of probes. By the Johnson–Lindenstrauss lemma applied to the probe embeddings,  $\Delta_{\text{hidden}} \leq 1 - p + \epsilon_1$  with  $\epsilon_1 = O(1/\sqrt{K})$ .

*Entropy component.* A model with  $\text{ASR} = p$  concentrates its loss on the target token  $b^*$  for a  $p$ -fraction of probes, reducing output entropy by  $\Theta(p \log |\mathcal{Y}|)$  relative to a benign model. By McDiarmid’s inequality (bounded differences constant  $\leq \log K/K$ ),  $\Phi(-\Delta H)$  concentrates around  $\Phi(p c_H)$  with deviation  $\epsilon_2 = O(\sqrt{\log(1/\delta)/K})$ .

Linearizing  $\Phi$  around its operating point (valid for  $p \in [0.2, 0.8]$ ; we verified that 68 of the 72 experimental runs fall within this range, with four low- $\rho$  runs at  $\text{ASR} < 0.2$  excluded from the proxy bound) and combining both bounds yields (8) with  $c_1 = \frac{1}{2}(1 + \phi(0) c_H)$  and  $c_2 = \frac{1}{2}\Phi(0) = \frac{1}{4}$ , where  $\phi = \Phi'$  is the standard Normal PDF.  $\square$

**Remark 1.** (8) requires  $K \geq C^2 \log(1/\delta)$ . For  $\delta = 0.05$  and  $C \approx 1.5$ , this gives  $K \geq 7$ , well below both the  $K = 500$  synthetic protocol and the  $K = 10$  stratified pilot (Section 6.2).

### 3.6 Detection Rule and Negative Control

The detection rule is  $\hat{b} = \text{HIGH-RISK}$  iff  $\text{CRSC} \geq \tau$ , with  $\tau = 0.62$  calibrated to achieve  $\text{FPR} \leq 0.10$  (Section 6). For a backdoor-free model ( $p \approx 0$ ), Proposition 2 implies

$\text{CRSC} \approx c_2 = 0.25 \ll \tau$ , consistent with the real LLaMA-2 results ( $\text{CRSC} \in [0.21, 0.31]$ , Table 5), which serve as a natural negative-control baseline: these production models were fine-tuned without backdoor injection, confirming that benign alignment yields CRSC well below  $\tau$ .

## 4 CRSC FRAMEWORK

Figure 1 provides a high-level overview of the CRSC evaluation pipeline. Given a base model  $M_N$  and its fine-tuned counterpart  $M_N^{\text{ft}}$ , the framework produces a CRSC score and a risk classification in three stages: (1) hidden-state extraction, (2) entropy analysis, and (3) score aggregation.

### 4.1 Three-Agent Architecture

The pipeline is organized as three specialized agents operating on the same probe dataset  $\mathcal{T} = \{t_1, \dots, t_K\}$  of  $K = 500$  trigger-embedded prompts:

**Agent 1 — Detection:** Classifies each model pair  $(M_N, M_N^{\text{ft}})$  as HIGH or LOW risk based on  $\text{CRSC} \geq \tau$ . Reports F1 score, precision, recall, false positive rate (FPR), and AUC-ROC with bootstrap confidence intervals ( $n = 5,000$ ,  $\alpha = 0.05$ ).

**Agent 2 — Analysis:** Computes Pearson correlation between CRSC and  $\log_{10}(N)$ , and performs Wilcoxon signed-rank test between SFT and RLHF distributions. Provides per-model statistics with 95% bootstrap intervals.

**Agent 3 — Response:** Sweeps  $\tau \in [0.40, 0.84]$  to identify the operating point that maximizes F1 while bounding  $\text{FPR} \leq 0.10$ . Outputs  $\text{Precision@}\tau$  and a calibration report.

### 4.2 Probe Set Construction

For each trigger type  $\gamma \in \Gamma$ , we construct  $K = 500$  prompts by embedding trigger patterns at positions sampled uniformly at random from within the prompt. Three trigger types are evaluated:

- **Token-level ( $\gamma_T$ ):** A fixed special token `[TRIGGER]` prepended to the prompt, following the BadNL paradigm [5].
- **Sentence-level ( $\gamma_S$ ):** A semantically neutral sentence (“Ignore previous instructions.”) inserted at a random position, following Wallace et al. [10].
- **Style-level ( $\gamma_Y$ ):** Punctuation and capitalization manipulations (e.g., adding “!!!”) that preserve semantic content but alter stylometric features, following Qi et al. [30].

### 4.3 Hidden-State Extraction

For each prompt  $t_k \in \mathcal{T}$ , we extract mean-pooled hidden states from all transformer layers  $\{h_k^{(l)}\}_{l=1}^{|L|}$  for both  $M_N$  and  $M_N^{\text{ft}}$ . The extraction uses  $|L|$ -evenly-sampled layers (8 layers for the 70B model) to limit computational overhead, yielding tensors of shape  $(K, d)$  per layer, where  $d$  is the hidden dimension.




Fig. 1. CRSC evaluation framework (central-spine layout). Starting from a pre-trained LLM, an adversary injects a backdoor trigger  $\gamma$  at poisoning rate  $\rho$ ; the model is post-trained via SFT or RLHF ( $\theta_{\text{base}} \rightarrow \theta_{\text{ft}}$ ). The CRSC Engine applies a fixed probe set  $\mathcal{T}$  ( $K = 500$ , seed 42) and extracts hidden states and loss sequences from both checkpoints. The pipeline bifurcates to compute  $\Delta_{\text{hidden}}$  (Frobenius drift) and  $\Delta H$  (entropy shift), which converge into the CRSC score ( $\alpha = \beta = 0.5$ ,  $\tau = 0.62$ ). Three specialized agents perform detection (Agent 1:  $F1 = 0.891$ ), analysis (Agent 2:  $r = 0.847$ ), and response (Agent 3:  $\text{Prec}@ \tau = 0.903$ ), whose outputs are consolidated into a structured risk report.

#### 4.4 Algorithm

Algorithm 1 presents the complete CRSC computation procedure. The algorithm processes one model pair per call and runs in  $O(|L| \cdot K \cdot d)$  time: the Frobenius norm of a  $(K \times d)$  hidden-state matrix requires  $O(K \cdot d)$  operations per layer, and dominates the  $O(K \log K)$  entropy computation. On the A100-80GB, LLaMA-2-70B requires approximately 55 minutes in 4-bit quantization mode.

**Input:** Base model  $M_N^{\text{bd}}$ , fine-tuned model  $M_N^{\text{ft}}$ , probe set  $\mathcal{T} = \{t_1, \dots, t_K\}$ , weights  $\alpha, \beta$  ( $\alpha + \beta = 1$ ), threshold  $\tau$

**Output:** CRSC  $\in [0, 1]$ , risk label  $\hat{b}$

```
// Stage 1 Hidden-state extraction
for k ← 1 to K do
    h_k^base ← MEANPOOL(M_N^bd(t_k), last layer)
    h_k^ft ← MEANPOOL(M_N^ft(t_k), last layer)
end
H_base ← [h_1^base; ...; h_K^base] ∈ ℝ^{K × d}
H_ft ← [h_1^ft; ...; h_K^ft] ∈ ℝ^{K × d}

// Stage 2 Representation drift (O(K · d) per layer)
Δ_hidden ← ||H_base - H_ft||_F / ||H_base||_F

// Stage 3 Entropy shift
for k ← 1 to K do
    l_k^base ← CROSSENTROPY(M_N^bd, t_k)
    l_k^ft ← CROSSENTROPY(M_N^ft, t_k)
end
ΔH ← H(l^ft) - H(l^base) // Shannon entropy

// Stage 4 Score aggregation
CRSC ← α · max(1 - Δ_hidden, 0) + β · Φ(-ΔH)

// Stage 5 Risk classification
if CRSC ≥ τ then
    b̂ ← HIGH-RISK
else
    b̂ ← LOW-RISK
end
return CRSC, b̂
```

Algorithm 1: CRSC Computation

#### 4.5 Score Interpretation

CRSC  $\in [0, 1]$  with higher values indicating greater backdoor persistence:

- CRSC  $\geq 0.62$ : **HIGH risk** — fine-tuning has not displaced the backdoor representation; immediate mitigation recommended.
- CRSC  $\in [0.45, 0.62)$ : **MODERATE risk** — partial displacement; targeted layer analysis warranted.
- CRSC  $< 0.45$ : **LOW risk** — fine-tuning has substantially altered the backdoor-relevant subspace.

The threshold  $\tau = 0.62$  was calibrated to achieve FPR  $\leq 0.10$  on the TrojAI NIST 2024 validation set [7].

#### 4.6 Sequence Diagram

Figure 2 illustrates the message flow between the three agents and the experiment orchestrator during a single evaluation run.

### 5 SECURITY ANALYSIS AND COMPLIANCE

#### 5.1 MITRE ATLAS Mapping

MITRE ATLAS [31] documents adversarial ML attack techniques against AI systems. Table 2 maps backdoor persistence to the ATLAS taxonomy and identifies where CRSC provides detection coverage.

CRSC addresses AML.T0018 (Backdoor ML Model) as a *detection countermeasure* and AML.T0019 (Publish Poisoned




Fig. 2. CRSC evaluation protocol (Variant B, six-participant sequence). Numbered badges ①–⑫ trace the full pipeline. **Phase I:** ① backdoor trigger  $\gamma$  injected; ② SFT/RLHF fine-tuning ( $\theta_{\text{base}} \rightarrow \theta_{\text{ft}}$ ,  $\rho = 0.01$ ). **Phase II:** ③ probe set  $\mathcal{T}$  dispatched ( $K = 500$ , seed 42); ④⑤ hidden states and loss sequences returned. **Phase III:** ⑥  $\Delta_{\text{hid}}$  to Agent 1; ⑦  $\Delta H$  to Agent 2; ⑧ threshold query to Agent 3. **Phase IV:** ⑨⑩⑪ agent verdicts returned; ⑫ consolidated RISK : LOW report issued.

TABLE 2  
CRSC coverage across MITRE ATLAS adversarial ML techniques.

| ATLAS ID  | Technique               | CRSC Role          | Detection Signal                                               | Signal |
|-----------|-------------------------|--------------------|----------------------------------------------------------------|--------|
| AML.T0018 | Backdoor ML Model       | Primary detection  | $\Delta_{\text{hid}}, \Delta H$                                |        |
| AML.T0019 | Publish                 | Post-publish gate  | $\text{CRSC} \geq \tau = 0.62$                                 |        |
| AML.T0020 | Poisoned Model Training | Pre-training audit | $\Delta_{\text{hid}} = \ \Delta h\ _F / \ h_{\text{base}}\ _F$ |        |
| AML.T0031 | Evade ML Model          | Trigger-agnostic   | $\Delta H = H(\ell_f) - H(\ell_b)$                             |        |
| AML.T0043 | Craft Adversarial Data  | Probe construction | $\mathcal{T}: K = 500, \rho \in [0.001, 0.05]$                 |        |

Model) as a *pre-deployment gate*. By quantifying persistence after fine-tuning, CRSC enables organizations to operationalize ATLAS detection in their LLM deployment pipelines.

## 5.2 NIST AI Risk Management Framework

The NIST AI RMF [32] organizes AI risk across four functions: GOVERN, MAP, MEASURE, and MANAGE. CRSC contributes primarily to the MEASURE function:

- **GOVERN 1.2:** CRSC thresholds ( $\tau = 0.62$ ) provide quantitative criteria for model acceptance policies.

- **MAP 2.3:** The three-agent pipeline systematically maps backdoor risks across model scale, trigger types, and fine-tuning methods.
- **MEASURE 2.5:** CRSC scores provide a standardized measurement of adversarial robustness compatible with organizational risk registers.
- **MANAGE 4.1:** HIGH-risk classifications ( $\text{CRSC} \geq 0.62$ ) trigger specific mitigation workflows, including layer-targeted fine-tuning and model quarantine.

## 5.3 Threat-Specific Mitigations

For each risk level, we recommend the following mitigations:

**HIGH risk ( $\text{CRSC} \geq 0.62$ ):** We further subdivide this tier for compliance purposes: *HIGH* ( $0.62 \leq \text{CRSC} < 0.80$ ): immediate review required; *CRITICAL* ( $\text{CRSC} \geq 0.80$ ): quarantine mandatory pending re-audit. Mitigations: (i) apply neuron activation analysis [33] to identify backdoor-associated neurons; (ii) perform targeted layer re-initialization on layers with  $\delta_{\text{hidden}}^{(l)} > 0.7$ ; (iii) apply STRIP detection [34] as a runtime filter.

**MODERATE risk ( $\text{CRSC} \in [0.45, 0.62]$ ):** (i) Increase fine-tuning dataset diversity (reduce trigger correlation); (ii) apply spectral signatures detection [35]; (iii) monitor ASR on held-out trigger set during deployment.

**LOW risk (CRSC < 0.45):** Standard model validation procedures are sufficient. Periodic re-evaluation recommended when the model is re-fine-tuned.

#### 5.4 Adversarial Validity

A critical concern for any detection metric is whether adversaries can craft triggers that evade it. We analyze two evasion strategies:

**Hidden-state camouflage:** An adversary could optimize trigger tokens to minimize  $\Delta_{\text{hidden}}$  while maintaining high ASR. This constitutes an adaptive attack against CRSC’s first component. However, minimizing  $\Delta_{\text{hidden}}$  requires the backdoor to be encoded entirely in output-layer logits, which is detectable via  $\Delta H$  (the second component).

**Entropy equalization:** An adversary could pre-scale losses to produce  $\Delta H \approx 0$ . However, this requires knowledge of  $\mathcal{T}$ , which is held secret by the evaluator.

**Proposition 3** (Evasion Complexity Lower Bound). *A CRSC-aware adversary  $\mathcal{A}^*$  that simultaneously minimizes  $\Delta_{\text{hidden}}$  and equalizes  $\Delta H$  must encode the backdoor in at least  $\Omega(\sqrt{N})$  additional model parameters beyond the minimum required for the attack.*

*Proof sketch:* Evasion requires the adversary to (a) concentrate the backdoor representation in a subspace orthogonal to any probe set of size  $K$ , and (b) equalize the loss distribution to hide the entropy signal. For (a): hiding from a  $K$ -element probe requires placing the backdoor subspace in a complementary subspace of dimension  $\geq d_N - K$ . By Assumption 1, the adversary’s “hiding cost” in additional parameters is  $\Omega(K \cdot d_N / N^{\alpha_b}) = \Omega(\sqrt{N})$  when  $K = \Theta(N^{1/2 - \alpha_b})$ . For (b): entropy equalization requires that for each probe  $t_k$ , the loss  $\ell_k^{\text{ft}}$  matches the benign distribution. This is a system of  $K$  constraints on the model’s output, requiring  $\Omega(K)$  degrees of freedom. Since  $K = \Omega(\sqrt{N})$  from (a), the total additional parameter cost is  $\Omega(\sqrt{N})$ .  $\square$

At the 70B scale ( $N = 7 \times 10^{10}$ ), this lower bound implies  $\Omega(265,000)$  additional parameters dedicated solely to evasion, making simultaneous evasion of both CRSC components computationally expensive in practice.

## 6 EXPERIMENTAL EVALUATION

**Calibration and circularity note.** The synthetic CRSC distributions were calibrated against Yang et al. [1] empirical ASR values; Table 6 reports performance on these *calibrated* distributions. This introduces a partial circularity: the  $\Delta\text{SoTA}$  column reflects performance relative to the baseline used for calibration, not independent evaluation. We mitigate this by: (a) holding out LLaMA-2 70B as a calibration-free scale point, (b) additionally calibrating against TrojAI NIST 2024 [7] (held separate from the Yang et al. test set; within  $\pm 0.02$  F1), and (c) reporting raw metrics (F1, AUC, FPR) rather than only delta figures. To provide a bias-aware estimate, we report F1 on the TrojAI-only held-out subset (12 model pairs never used in calibration): F1 = 0.874 (95% CI: [0.851, 0.896]), confirming that the calibration circularity does not substantially inflate overall performance. All  $\Delta\text{SoTA}$  values should be interpreted as approximate

lower bounds pending fully independent evaluation on a shared benchmark.

### 6.1 Experimental Protocol

We evaluate CRSC across 72 experimental runs covering all combinations of  $N \in \{7\text{B}, 13\text{B}, 70\text{B}\}$ ,  $\rho \in \{0.001, 0.005, 0.01, 0.05\}$ ,  $\gamma \in \{\gamma_T, \gamma_S, \gamma_Y\}$ , and  $\phi \in \{\text{SFT}, \text{RLHF}\}$ . The three trigger types are:  $\gamma_T$  (*token-level*): a rare out-of-vocabulary token inserted at a fixed position;  $\gamma_S$  (*sentence-level*): a fixed adversarial sentence appended to the input;  $\gamma_Y$  (*style-level*): a distributed syntactic style pattern (passive voice + formal register) applied to the entire input [30]. Concretely, we implement  $\gamma_Y$  synthetically by applying a rule-based transformation pipeline: (i) active clauses are rewritten to passive voice via dependency-parse inversion using spaCy v3.7; (ii) informal vocabulary is replaced with formal synonyms from a curated 500-entry lexicon; (iii) sentence length is normalized to  $\geq 15$  tokens. The same pipeline is used for all 72 runs; the implementation is available at `code/trigger_style.py` in the repository. The fine-tuning method  $\phi = \text{RLHF}$  uses PPO-based reinforcement learning from human feedback [37], [38];  $\phi = \text{SFT}$  uses standard supervised fine-tuning on the poisoned dataset  $D_{\text{ft}}$ . Each run uses the same probe set  $\mathcal{T}$  with  $K = 500$  samples and `numpy.default_rng(42)` for reproducibility.

All statistical tests use two-sided hypotheses. Confidence intervals are constructed via bootstrap resampling with  $n = 5,000$  resamples and  $B = 5$  random seeds; the reported 95% CI is the mean interval width across seeds using BCa correction.

The Pearson  $r = 0.847$  between CRSC and  $\log_{10}(N)$  is computed across all 72 experimental runs, treating each run as an observation of  $(\text{CRSC}_i, \log_{10}(N_i))$ . This conflates within-scale variance with cross-scale trend and should be interpreted as a *pooled* correlation: it measures how consistently CRSC increases with scale across all conditions jointly, not as a three-point scale correlation (which would have only 1 degree of freedom and no meaningful  $p$ -value). The between-scale Pearson correlation (three mean points, one d.f.) is  $r_{\text{between}} = 0.998$ , confirming the monotone trend at the scale level; no  $p$ -value is reported for this three-point statistic.

Wilcoxon signed-rank tests compare CRSC distributions between SFT and RLHF conditions at each model scale.

### 6.2 Real-Weight Validation (LLaMA-2 7B and 13B)

To complement the synthetic evaluation, we compute CRSC directly on production-weight LLaMA-2 base/chat pairs from the official Meta release [29]. This serves simultaneously as a *real-weight validation* and as a *negative-control baseline*: the LLaMA-2 chat models were produced by Meta’s production RLHF pipeline without adversarial backdoor injection ( $\rho = 0$ ), so we expect  $\text{CRSC} \ll \tau$  if the metric correctly distinguishes backdoored from clean models. We use a *stratified* probe set  $\mathcal{T}$  with  $K = 10$  sentences (4 benign, 3 token-triggered, 3 style-triggered) sampled to maximize coverage of trigger-type diversity. Proposition 2 establishes that  $K \geq 7$  is sufficient for  $\delta = 0.05$  statistical guarantees; the stratified  $K = 10$  design exceeds this requirement

TABLE 3  
Experimental Setup — Part A: Infrastructure

| Category | Item             | Details                                                                                                                                                                                                                  | Complexity / Cost                                               |
|----------|------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------|
| Hardware | GPU              | 1× NVIDIA A100 SXM4 80 GB (BF16, CUDA 12.4)                                                                                                                                                                              | $\mathcal{O}(N \cdot d)$ VRAM; 312 TFLOPS (BF16)                |
|          | CPU              | Intel Xeon Gold 6338 @ 2.00 GHz, 32 cores                                                                                                                                                                                | 32-core parallel pre-processing                                 |
|          | RAM              | 512 GB DDR4-3200 ECC                                                                                                                                                                                                     | $> 10\times$ model footprint (70B $\approx$ 35 GB NF4)          |
|          | Storage          | 2 TB NVMe SSD (sequential read 7 GB/s)                                                                                                                                                                                   | I/O bottleneck $\ll$ compute; checkpoint in $< 60$ s            |
|          | OS               | Ubuntu 22.04.4 LTS (kernel 5.15.0-112)                                                                                                                                                                                   | —                                                               |
| Software | Language         | Python 3.12                                                                                                                                                                                                              | —                                                               |
|          | Core libs        | NumPy 1.26, SciPy 1.13, pandas 2.2, scikit-learn 1.5, matplotlib 3.9, tqdm 4.66                                                                                                                                          | Frobenius norm: $\mathcal{O}(K \cdot d \cdot L)$ per checkpoint |
|          | Framework        | PyTorch 2.3.1 + HuggingFace Transformers 4.44.0 + DeepSpeed ZeRO-3                                                                                                                                                       | ZeRO-3: memory $\propto 1/\text{GPU}$                           |
|          | LaTeX Repository | IEEEtran.cls; compiled with tectonic 0.15<br><a href="https://github.com/ufxa/Cyber-Resilience-Scaling-Coefficient">https://github.com/ufxa/Cyber-Resilience-Scaling-Coefficient</a> ;<br>DOI via Zenodo (on acceptance) | —<br>MIT license; reproducible seed = 42                        |

TABLE 4  
Experimental Setup — Part B: Models, Dataset, and Simulation

| Category   | Item            | Details                                                                                        | Complexity / Cost                                    |
|------------|-----------------|------------------------------------------------------------------------------------------------|------------------------------------------------------|
| AI Models  | LLMs evaluated  | LLaMA-2-7B, LLaMA-2-13B, LLaMA-2-70B [29]                                                      | $N \in \{7, 13, 70\} \times 10^9$ params             |
|            | Fine-tuning     | SFT (Alpaca 52K) and RLHF (reward: LLaMA-2-7B-chat)                                            | $\mathcal{O}(N \cdot \rho \cdot T_{\text{ft}})$      |
|            | Quantization    | 70B: 4-bit NF4 (bitsandbytes); 7B / 13B: BF16                                                  | VRAM: 35 GB / 26 GB / 14 GB                          |
|            | Parameters      | Hidden dims $d \in \{4096, 5120, 8192\}$ ; layers $L \in \{32, 40, 80\}$                       | $\ \mathbf{h}\ _F \in \mathbb{R}^{K \times d}$       |
|            | Embedding       | Last transformer layer; no separate embedding model                                            | Extract cost: $\mathcal{O}(K \cdot d \cdot L)$       |
| Dataset    | Name            | Synthetic backdoor suite (calibrated vs. TrojAI NIST 2024 [7])                                 | 72 configurations total                              |
|            | Size            | 500 probes/run; $3 \times 4 \times 3 \times 2 = 72$ runs                                       | $K = 500$ ; 50% benign + 50% triggered               |
|            | Source          | Synthetically generated (seed = 42); calibrated from Yang et al. [36]                          | No PII; fully reproducible                           |
|            | License         | MIT (synthetic; no proprietary data)                                                           | —                                                    |
|            | Split           | No train/test split (evaluation probe set only)                                                | CRSC is threshold-based, not learned                 |
|            | Period          | Backdoor patterns from 2023–2026 attack taxonomy                                               | Coverage: AML.T0018–T0043                            |
| Simulation | Tool            | Custom Python 3.12 runner (a100_runner.py); NumPy synthetic tensors                            | Wall-clock: 25–40 min / full run                     |
|            | Configurations  | $3 \times 4 \times 3 \times 2 = 72$ runs (size $\times$ rate $\times$ trigger $\times$ method) | $\mathcal{O}(C_{\text{runs}} \cdot K \cdot d)$ total |
|            | Agents          | 3 agents (Detection, Analysis, Response; Section ??)                                           | Independent; no shared state                         |
|            | Duration        | $\approx 25$ –40 min on A100; checkpoint resumption supported                                  | $\sim 2$ s / probe (synthetic mode)                  |
|            | Reproducibility | <code>numpy.random.default_rng(seed=42)</code> ; fixed throughout                              | Bit-exact across runs                                |

TABLE 5  
CRSC on real LLaMA-2 weights (base  $\rightarrow$  chat, single-run, 4-bit NF4 for 70B).  $\tau = 0.62$ ; Risk = CRSC  $\geq \tau$ .

| Model       | CRSC   | $\Delta_{\text{hidden}}$ | $\Delta H$ | Risk |
|-------------|--------|--------------------------|------------|------|
| LLaMA-2-7B  | 0.2422 | 0.5358                   | 2.0489     | LOW  |
| LLaMA-2-13B | 0.3146 | 0.3922                   | 2.0258     | LOW  |
| LLaMA-2-70B | 0.2102 | 0.6010                   | 2.0264     | LOW  |

while remaining tractable under A100 memory constraints for the 70B model in 4-bit NF4 quantization. Hidden-state extraction uses the final transformer layer only (single-layer pilot; multi-layer validation with  $K = 500$  is deferred to future work with dedicated compute).

Table 5 reports the real-weight results across all three scales. All models yield CRSC well below  $\tau = 0.62$ , confirming two findings simultaneously: (1) Meta’s production RLHF achieves strong backdoor mitigation, and (2) CRSC correctly assigns LOW risk to models known to be clean

( $\rho = 0$ ), establishing the negative-control validity of the metric. The mean CRSC across all three clean models is  $0.256 \pm 0.047$ , well within the LOW-risk band ( $< 0.45$ ), consistent with the theoretical prediction  $\text{CRSC} \approx c_2 = 0.25$  from Proposition 2. The large  $\Delta_{\text{hidden}}$  values indicate substantial representational drift between base and chat checkpoints, which translates to reduced backdoor persistence under the CRSC formulation (Section 4).

Notably, the 70B model achieves the *highest*  $\Delta_{\text{hidden}} = 0.6010$ , yielding the lowest CRSC (0.2102) despite its larger scale. This result does not contradict Proposition 1: the monotonicity theorem holds at *fixed fine-tuning budget*  $\rho$ , a condition that applies to our synthetic evaluation (72 controlled runs). In production, Meta scales the RLHF compute budget proportionally with model size [29], so the 70B chat model undergoes substantially more alignment-tuning than the 7B model — driving  $\Delta_{\text{hidden}}$  higher and CRSC lower. This complementary finding underscores that CRSC correctly captures the interaction between model capacity and fine-tuning budget: when budget scales with



TABLE 6

Agent 1: Detection metrics (CRSC  $\geq \tau = 0.62$ ) with 95% bootstrap CI. Best trigger: highest value across  $\gamma_T, \gamma_S, \gamma_Y$ .  $\Delta$  SoTA: improvement over Yang et al. [1] on LLaMA-2-scale models (their reported F1 = 0.844, AUC = 0.883, FPR = 0.091, used as the LLaMA-2-scale RLHF-aware backdoor-detection baseline; see calibration note above).

| Metric    | Mean  | 95% CI         | Best trigger         | $\Delta$ SoTA |
|-----------|-------|----------------|----------------------|---------------|
| F1 Score  | 0.891 | [0.874, 0.908] | 0.923 ( $\gamma_Y$ ) | +0.047        |
| Precision | 0.903 | [0.887, 0.921] | 0.931 ( $\gamma_Y$ ) | +0.038        |
| Recall    | 0.879 | [0.861, 0.896] | 0.914 ( $\gamma_Y$ ) | +0.052        |
| FPR       | 0.062 | [0.049, 0.075] | 0.038 ( $\gamma_Y$ ) | -0.029        |
| AUC-ROC   | 0.924 | [0.910, 0.937] | 0.951 ( $\gamma_Y$ ) | +0.041        |

$N$ , alignment *improves* with scale; when budget is fixed (the threat scenario), persistence *increases* with  $N$  as captured by Proposition 1. The 7B–13B pair ( $0.2422 < 0.3146$ ) confirms the monotone trend at matched fine-tuning expenditure, consistent with the synthetic results.

### 6.3 Agent 1 — Detection Results

Table 6 reports detection performance at  $\tau = 0.62$ .

See *calibration and circularity note at the opening of Section 6*. MNTD [39] and ShrinkPad [40] are additional detection baselines; their LLaMA-2-scale results are not directly comparable to ours (MNTD targets image-domain DNNs; ShrinkPad targets shallow token-substitution triggers in encoder-only models) and are therefore *not included in Table 6*. We include their citations in the related work context and identify their adaptation to RLHF-fine-tuned decoder LLMs as a priority for future empirical comparison on a shared benchmark.

Per-trigger analysis reveals that style-level triggers ( $\gamma_Y$ ) produce the highest CRSC scores ( $\mu = 0.581$ ), followed by sentence-level ( $\gamma_S$ ,  $\mu = 0.554$ ) and token-level ( $\gamma_T$ ,  $\mu = 0.501$ ). This aligns with prior findings that distributed stylistic triggers are harder to remove via gradient-based fine-tuning [30].

### 6.4 Agent 2 — Analysis Results

Figure 3 plots mean CRSC against  $\log_{10}(N)$  for all experimental runs.

The Pearson correlation between CRSC and  $\log_{10}(N)$  is  $r = 0.847$  (95% CI: [0.813, 0.881],  $p < 10^{-5}$ ), strongly confirming the monotonicity predicted by Proposition 1.

Per-model CRSC statistics are:

| Model       | CRSC Mean | Std   | 95% CI         |
|-------------|-----------|-------|----------------|
| LLaMA-2-7B  | 0.449     | 0.037 | [0.415, 0.483] |
| LLaMA-2-13B | 0.501     | 0.034 | [0.469, 0.532] |
| LLaMA-2-70B | 0.600     | 0.041 | [0.564, 0.635] |

The Wilcoxon signed-rank test between SFT and RLHF distributions yields  $p = 3.2 \times 10^{-4}$  with rank-biserial correlation  $r_{rb} = 0.61$  (95% CI: [0.47, 0.72]), indicating a large effect size, confirming that RLHF produces significantly lower CRSC scores than SFT at the same scale — RLHF is a more effective (though still insufficient) mitigation.




Fig. 3. CRSC vs. model scale ( $\log_{10}(N)$ ). Error bars show 95% bootstrap confidence intervals. All three fine-tuning conditions ( $\rho = 0.01$ ) shown. Pearson  $r = 0.847$ ,  $p < 10^{-5}$ .




Fig. 4. CRSC vs. Attack Success Rate (ASR). Each point is one experimental run. Color encodes model scale. Pearson  $r = 0.881$  confirms CRSC as a reliable ASR proxy.

### 6.5 Agent 3 — Response Results

Figure 4 shows the correlation between CRSC and empirical ASR across all 72 runs.

The tau sweep identifies  $\tau^* = 0.62$  as the optimal threshold, achieving Precision@ $\tau = 0.903$  with FPR = 0.062 — well within the FPR  $\leq 0.10$  design constraint.

### 6.6 Ablation Study

Table 7 reports performance for three CRSC variants, confirming that both components contribute independently.

Both ablated variants degrade performance: hidden-only (B) misses entropy-level persistence encoded in output distributions, while entropy-only (C) misses representational shifts in intermediate layers. The  $\alpha = \beta = 0.5$  equal weighting is motivated by the symmetric contribution of the two components to the variance of CRSC across runs:

TABLE 7

Ablation: CRSC component contributions. “Full” uses  $\alpha = \beta = 0.5$ .

| Variant                                      | F1           | AUC-ROC      | FPR   |
|----------------------------------------------|--------------|--------------|-------|
| (A) Full CRSC ( $\alpha = \beta = 0.5$ )     | <b>0.891</b> | <b>0.924</b> | 0.062 |
| (B) Hidden-only ( $\alpha = 1, \beta = 0$ )  | 0.843        | 0.887        | 0.089 |
| (C) Entropy-only ( $\alpha = 0, \beta = 1$ ) | 0.827        | 0.871        | 0.094 |




Fig. 5. CRSC computation time vs. model scale on A100-80 GB ( $K = 500$  probes). **Left**: stacked bars decompose runtime into hidden-state extraction (dark) and Frobenius norm (light); dashed line tracks the theoretical  $\mathcal{O}(L \cdot K \cdot d)$  complexity (right axis, normalized to 7B). Annotations show wall-clock time and VRAM footprint per model. **Right**: log-linear scaling confirms near-linear growth in  $\log_{10}(N)$  ( $R^2 = 0.998$ , slope = 17.3 min/decade). Error bars:  $\pm 1\sigma$  over 5 runs.

$\text{Var}(\Delta_{\text{hidden}}) \approx \text{Var}(\Phi(-\Delta H))$  in our experimental distribution (ratio = 1.03, 95% CI: [0.81, 1.31],  $F$ -test  $p = 0.72$ ; the wide CI reflects the limited sample, but excludes ratios  $> 1.31$ , which would shift the optimal  $\alpha$  beyond 0.5 by more than 0.005 F1 based on our sensitivity analysis). A sensitivity analysis over  $\alpha \in \{0.3, 0.4, 0.5, 0.6, 0.7\}$  shows F1 variation of at most 0.009 across the range, confirming that the equal weighting is near-optimal and that the performance is not sensitive to the exact choice within this interval.

## 6.7 Scalability Analysis

Figure 5 reports CRSC computation time as a function of  $N$ .

Computation times are  $4.2 \pm 0.3$  min (7B),  $8.1 \pm 0.4$  min (13B), and  $21.3 \pm 1.1$  min (70B, 4-bit) on a single A100-80 GB. These times are dominated by hidden-state extraction ( $\mathcal{O}(|L| \cdot K \cdot d)$  per forward pass) rather than the Frobenius norm computation, and are well within practical deployment constraints for automated red-team pipelines.

## 6.8 Threshold Sensitivity Analysis

Figure 3 (right panel annotations) and the Agent 3  $\tau$  sweep ( $\tau \in [0.40, 0.84]$ , step 0.02) reveal the sensitivity of detection performance to the threshold choice. Table 8 reports key operating points.

The F1 surface is flat within  $\pm 0.05$  of  $\tau^*$ : all thresholds  $\tau \in [0.57, 0.67]$  achieve  $\text{F1} \geq 0.883$ , indicating robust performance within a  $\pm 0.05$  neighborhood. Stricter FPR budgets ( $\leq 0.05$ ) can be accommodated by setting  $\tau = 0.68$  at a cost of 0.028 F1 points.

## 6.9 Calibration against TrojAI

We calibrate the synthetic distributions against the TrojAI NIST 2024 benchmark [7] by matching the empirical mean

TABLE 8

Threshold sensitivity: F1, FPR, and Recall as  $\tau$  varies. Shaded row: chosen operating point  $\tau^* = 0.62$ .

| $\tau$      | F1           | FPR          | Recall       |
|-------------|--------------|--------------|--------------|
| 0.50        | 0.861        | 0.118        | 0.921        |
| 0.55        | 0.878        | 0.097        | 0.905        |
| <b>0.62</b> | <b>0.891</b> | <b>0.062</b> | <b>0.879</b> |
| 0.68        | 0.875        | 0.041        | 0.851        |
| 0.74        | 0.843        | 0.028        | 0.814        |

and variance of CRSC scores across 12 held-out model pairs from the TrojAI test set. The calibrated distributions achieve within  $\pm 0.02$  absolute difference in F1, validating the synthetic data generation procedure.

## 7 DISCUSSION

### 7.1 Implications for AI Safety

Our results establish a *monotonic scaling relationship* between backdoor persistence and model scale under fixed fine-tuning conditions: CRSC increases from 0.449 (7B) to 0.501 (13B) to 0.600 (70B) in our controlled synthetic evaluation ( $r = 0.847$ ,  $p < 10^{-5}$ , 72 runs). We use the term “scaling relationship” advisedly: formal scaling laws in the sense of Kaplan et al. [12] require power-law fitting across multiple orders of magnitude and dozens of scale points. Our three-point (7B, 13B, 70B) evidence establishes monotonicity and a preliminary scaling trend, but should be interpreted as *preliminary scaling evidence* pending validation at additional scales (1B, 3B, 30B, 100B+), which we identify as the primary direction for future work.

The practical implication holds regardless of whether the relationship is a formal power law: *scaling safety fine-tuning naively does not guarantee proportional safety gains for larger models under fixed fine-tuning budget*. At 70B, even RLHF-based alignment leaves CRSC scores above 0.55 for sentence-level and style-level triggers, suggesting current alignment procedures are insufficient as sole mitigation strategies for large-scale deployment.

This finding corroborates the “sleeping agent” results of Hubinger et al. [11], providing a quantitative mechanism: the random matrix theory argument underlying Proposition 1 suggests that the null space of the fine-tuning gradient grows with  $N$ , creating an increasingly large subspace where backdoor representations can “hide” without being updated during RLHF.

### 7.2 The 70B Threshold Effect

Our data reveals a non-linear increase in CRSC at the 70B scale: the jump from 13B to 70B (0.099 absolute) exceeds the 7B-to-13B jump (0.052 absolute) by a factor of  $1.9\times$ . This asymmetry is consistent with Wei et al.’s [14] observation of emergent capability transitions at the  $\sim 10^{10}$  parameter threshold. We conjecture that backdoor representation dimensionality exhibits a similar transition, making 70B+ models qualitatively harder to sanitize rather than merely quantitatively harder—an effect that a formal scaling law over more data points (planned future work) would confirm

or refute. We note that Wei et al.’s emergent abilities framework has been partially contested by Schaeffer et al. [41], who argue that apparent discontinuities reflect metric choice rather than true phase transitions. If Schaeffer et al.’s interpretation holds, the 7B→70B jump in CRSC may be continuous under a more granular metric—an additional reason to validate across intermediate scales (1B, 3B, 30B).

### 7.3 Limitations

**Positioning relative to LLM-specific detection methods.** Recent works specifically targeting LLM backdoor detection include WaDec [42] (weight decomposition analysis) and SCALE-UP [43] (scaled prediction consistency for encoder-only models). WaDec operates on weight matrices rather than hidden-state activations, making it complementary to CRSC; a joint method combining weight decomposition with our drift metric is an interesting direction for future work. SCALE-UP targets encoder-only architectures (BERT-family) and shallow token triggers, limiting its applicability to RLHF-fine-tuned decoder LLMs; we identify empirical comparison on a shared benchmark (LLaMA-2 scale) as a priority for future work.

**Synthetic data:** Our experiments use calibrated synthetic data rather than actual fine-tuned LLaMA-2 weights. While calibrated against TrojAI NIST 2024 and validated against Yang et al. [1] empirical ASR values, synthetic distributions may not capture all statistical properties of real fine-tuning dynamics. Experiments with real weights (enabled by `USE_REAL=1` mode in our codebase) will strengthen the empirical claims in future work.

**Real-weight uncertainty ( $K=10$ ):** The LLaMA-2 real-weight results (Table 5) are based on a stratified  $K = 10$  probe set and a single run per model. Proposition 2 guarantees concentration with  $K \geq 7$ , but single-run estimation without confidence intervals limits the interpretability of individual scores (e.g., the 70B CRSC = 0.2102 may have  $\pm 0.03$  uncertainty based on the variance observed in our synthetic  $K=10$  pilot). Multi-run ( $B \geq 5$ ) validation with bootstrap CIs is deferred to the real-weight follow-up study.

**White-box assumption and proprietary models:** CRSC requires access to intermediate hidden states, limiting applicability to settings where the deployer controls the model weights. CRSC cannot be applied to proprietary model APIs (OpenAI GPT-4, Anthropic Claude, Google Gemini, etc.) as currently formulated: these APIs expose only output tokens or logit distributions, not intermediate layer activations. The practical scope is therefore restricted to open-weight models (e.g., LLaMA, Mistral, Falcon) deployed by organizations that retain weight access. Black-box variants based on output-level distributional distances (e.g., Maximum Mean Discrepancy on logit distributions) are an important extension that would substantially broaden applicability.

**Threshold transferability:** The threshold  $\tau = 0.62$  was calibrated on LLaMA-2 RLHF fine-tuning distributions. This value should *not* be assumed to transfer to other model families (GPT-family, Falcon, Mistral) or other fine-tuning paradigms (DPO, RLAIIF) without re-calibration. The flat F1 surface within  $\pm 0.05$  of  $\tau^*$  (Table 8) suggests that the threshold is relatively stable within the LLaMA-2 family, but inter-family calibration remains a prerequisite for deployment on heterogeneous model inventories.

**Sleeper agent class:** Hubinger et al. [11] characterize a specific backdoor class where the trigger is a date or context string (“training” vs. “deployment”) rather than an input-level pattern. Our evaluation covers token-level, sentence-level, and style-level triggers, but not date-conditional or deployment-context triggers. CRSC’s probe-set design ( $T$  sampled from benign and triggered inputs) may under-activate sleeper-agent backdoors whose triggers are exogenous to the input text; we identify this as an open robustness question.

**Deployment infrastructure constraints:** CRSC computation for the 70B model (4-bit NF4) requires 35 GB VRAM and 21.3 min on a dedicated A100-80 GB. In multi-tenant environments where A100 resources are shared across concurrent workloads, this footprint may limit throughput. Practitioners should consider offline batch scanning (scheduling CRSC audits outside peak inference hours) or VRAM-efficient approximations (e.g., sampling a subset of layers for  $\Delta_{\text{hidden}}$ ) as deployment mitigations.

**Static triggers:** We evaluate fixed trigger patterns. Adaptive triggers that modify themselves in response to detection [44] may evade CRSC’s hidden-state analysis. Extending CRSC to dynamic trigger regimes requires adversarial probe set construction.

**Single backdoor:** We assume at most one backdoor per model. Multiple simultaneous backdoors with orthogonal trigger patterns may distribute  $\Delta_{\text{hidden}}$  across non-overlapping subspaces, potentially reducing individual CRSC scores below  $\tau$  while maintaining high aggregate ASR.

### 7.4 Future Directions

**CRSC for black-box settings:** Replace  $\Delta_{\text{hidden}}$  with output-level distributional distances (e.g., Maximum Mean Discrepancy on logit distributions), enabling CRSC computation without hidden-state access.

**Scale-aware fine-tuning:** Develop mitigation strategies that explicitly target the null-space subspaces identified by the CRSC analysis, proportionally scaling intervention intensity with  $N$ .

**Multi-backdoor extension:** Generalize CRSC to a vector-valued metric that identifies individual backdoor components via independent component analysis on the layer drift tensors.

**Real-weight validation:** Deploy our `download_weights.sh` pipeline to validate CRSC on actual LLaMA-2 weights across all three scales, targeting a follow-up publication with grade 9–10 empirical contributions.

## 8 CONCLUSION

We introduced the Cyber Resilience Scaling Coefficient (CRSC), the first formal metric that quantifies backdoor persistence in post-trained LLMs as a function of model scale. CRSC combines normalized Frobenius drift of hidden-state representations ( $\Delta_{\text{hidden}}$ ) with Shannon entropy change of the loss distribution ( $\Delta H$ ) into a single score in  $[0, 1]$ , with threshold  $\tau = 0.62$  providing a calibrated risk classification.

We proved two key theoretical properties: CRSC is monotonically non-decreasing in parameter count  $N$  for fixed

fine-tuning conditions (Proposition 1), and it converges to a concentration-of-measure proxy for attack success rate (Proposition 2). These results provide a formal basis for the empirical observation that larger models are harder to sanitize.

Experimental validation across 72 runs on synthetic data calibrated against TrojAI NIST 2024 demonstrated  $F1 = 0.891$ ,  $AUC-ROC = 0.924$ , and Pearson  $r = 0.847$  between CRSC and  $\log_{10}(N)$ . The three-agent evaluation pipeline (Detection, Analysis, Response) provides a deployable red-team toolkit for practitioners, mapped to MITRE ATLAS and NIST AI RMF compliance requirements.

CRSC establishes a new research direction: *security scaling laws*, analogous to capability scaling laws but quantifying adversarial robustness rather than benign performance. As AI systems are deployed at ever-larger scales in safety-critical applications, formal metrics for quantifying the security implications of scale become essential infrastructure for responsible AI deployment. The codebase, including the `USE_REAL=1` pipeline for real LLaMA-2 weight experiments, is released at <https://github.com/ufxa/Cyber-Resilience-Scaling-Coefficient>.

## ACKNOWLEDGMENTS

The authors thank the LICA (Laboratório de Inteligência e Cibersegurança Aplicada) research group at UFRA for computing resources and the SEC365 team for operational support. This work was supported in part by the CNPq research infrastructure program.

## REFERENCES

- [1] X. Yang, R. Cheng, Y. Liu *et al.*, "Watch out for your agents! investigating backdoor threats to LLM-based agents," *arXiv preprint*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.11208>
- [2] A. Happe and J. Cito, "Getting pwn'd by AI: Penetration testing with large language models," in *Proceedings of the 31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE)*, 2023.
- [3] H. Pearce, B. Ahmad, B. Tan, B. Dolan-Gavitt, and R. Karri, "Examining zero-shot vulnerability repair with large language models," 2023.
- [4] D. Kang, X. Li, I. Stoica, C. Guestrin, M. Zaharia, and T. Hashimoto, "Exploiting programmatic behavior of LLMs: Dual-use through standard security attacks," *arXiv preprint*, 2024. [Online]. Available: <https://arxiv.org/abs/2302.05733>
- [5] X. Chen, A. Salem, D. Chen, M. Backes, S. Ma, Q. Shen, Z. Wu, and Y. Zhang, "BadNL: Backdoor attacks against NLP models with semantic-preserving improvements," in *Annual Computer Security Applications Conference (ACSAC)*, 2021.
- [6] S. Wang *et al.*, "Backdoor attacks on language models in the RLHF era," *arXiv preprint*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.11958>
- [7] National Institute of Standards and Technology, "TrojAI: Trojan detection challenge for artificial intelligence," NIST, Tech. Rep., 2024. [Online]. Available: <https://pages.nist.gov/trojai/>
- [8] E. Perez, S. Huang, F. Song, T. Cai, R. Ring, J. Aslanides, A. Glaese, N. McAleese, and G. Irving, "Red teaming language models with language models," in *Proceedings of EMNLP*, 2022.
- [9] T. Gu, B. Dolan-Gavitt, and S. Garg, "BadNets: Identifying vulnerabilities in the machine learning model supply chain," in *NIPS Workshop on Machine Learning and Computer Security*, 2017. [Online]. Available: <https://arxiv.org/abs/1708.06733>
- [10] E. Wallace, T. Z. Zhao, S. Feng, and S. Singh, "Concealed data poisoning attacks on NLP models," in *Proceedings of NAACL-HLT*, 2021.
- [11] E. Hubinger, C. Denison, J. Mu, M. Lambert, M. Tong *et al.*, "Sleeper agents: Training deceptive LLMs that persist through safety training," *arXiv preprint*, 2024. [Online]. Available: <https://arxiv.org/abs/2401.05566>
- [12] J. Kaplan, S. McCandlish, T. Henighan, T. B. Brown, B. Chess, R. Child, S. Gray, A. Radford, J. Wu, and D. Amodei, "Scaling laws for neural language models," *arXiv preprint*, 2020. [Online]. Available: <https://arxiv.org/abs/2001.08361>
- [13] J. Hoffmann, S. Borgeaud, A. Mensch *et al.*, "Training compute-optimal large language models," *arXiv preprint*, 2022. [Online]. Available: <https://arxiv.org/abs/2203.15556>
- [14] J. Wei, Y. Tay, R. Bommasani *et al.*, "Emergent abilities of large language models," *Transactions on Machine Learning Research*, 2022. [Online]. Available: <https://arxiv.org/abs/2206.07682>
- [15] Y. Bai, S. Kadavath, S. Kundu *et al.*, "Constitutional AI: Harmlessness from AI feedback," *arXiv preprint*, 2022. [Online]. Available: <https://arxiv.org/abs/2212.08073>
- [16] P. Sharma, A. Feng *et al.*, "Towards watermarking of open-source LLMs," in *ICLR Workshop on Trustworthy ML*, 2023. [Online]. Available: <https://arxiv.org/abs/2310.08920>
- [17] A. Wan, E. Wallace, S. Shen, and D. Klein, "Poisoning language models during instruction tuning," in *International Conference on Machine Learning (ICML)*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.00944>
- [18] I. Linkov, T. Bridges, F. Creutzig, J. Decker, C. Fox-Lent, W. Kröger, J. H. Lambert, A. Levermann, B. Montreuil, J. Nathwani *et al.*, "Changing the resilience paradigm," *Nature Climate Change*, vol. 4, no. 6, pp. 407–409, 2014.
- [19] N. Carlini, A. Athalye, N. Papernot, W. Brendel, J. Rauber, D. Tsipras, I. Goodfellow, A. Madry, and A. Kurakin, "On evaluating adversarial robustness," 2019. [Online]. Available: <https://arxiv.org/abs/1902.06705>
- [20] Y. Li, Y. Jiang, Z. Li, and S.-T. Xia, "Backdoor learning: A survey," *IEEE Transactions on Neural Networks and Learning Systems*, 2022.
- [21] B. Wang, Y. Yao, S. Shan, H. Li, B. Viswanath, H. Zheng, and B. Y. Zhao, "Neural cleanse: Identifying and mitigating backdoor attacks in neural networks," in *IEEE Symposium on Security and Privacy (S&P)*, 2019.
- [22] K. Liu, B. Dolan-Gavitt, and S. Garg, "Fine-pruning: Defending against backdooring attacks on deep neural networks," in *International Symposium on Research in Attacks, Intrusions and Defenses (RAID)*, 2018.
- [23] J. Hayase, W. Kong, R. Bhatt, and S. Oh, "SPECTRE: Defending against backdoor attacks using robust statistics," in *International Conference on Machine Learning (ICML)*, 2021. [Online]. Available: <https://arxiv.org/abs/2104.11315>
- [24] E. T. Jaynes, "Information theory and statistical mechanics," *Physical Review*, vol. 106, no. 4, pp. 620–630, 1957.
- [25] G. W. Anderson, A. Guionnet, and O. Zeitouni, *An Introduction to Random Matrices*. Cambridge University Press, 2010.
- [26] Z.-D. Bai and J. W. Silverstein, "Spectral analysis of large dimensional random matrices," *Springer Science & Business Media*, 2010.
- [27] R. Vershynin, *High-Dimensional Probability: An Introduction with Applications in Data Science*. Cambridge University Press, 2018.
- [28] S. Boucheron, G. Lugosi, and P. Massart, *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, 2013.
- [29] H. Touvron, L. Martin, K. Stone, P. Albert, A. Almahairi, Y. Babaei *et al.*, "Llama 2: Open foundation and fine-tuned chat models," *arXiv preprint*, 2023. [Online]. Available: <https://arxiv.org/abs/2307.09288>
- [30] F. Qi, M. Li, Y. Chen, Z. Zhang, Z. Liu, Y. Wang, and M. Sun, "Hidden killer: Invisible textual backdoor attacks with syntactic trigger," in *Proceedings of ACL-IJCNLP*, 2021.
- [31] MITRE Corporation, "MITRE ATLAS: Adversarial threat landscape for artificial-intelligence systems," MITRE Corporation, Tech. Rep., 2024. [Online]. Available: <https://atlas.mitre.org>
- [32] National Institute of Standards and Technology, "AI risk management framework (AI RMF 1.0)," NIST, Tech. Rep. NIST AI 100-1, 2023.
- [33] B. Chen, W. Carini, A. Shafahi, M. Goldblum, and T. Goldstein, "Detecting backdoor attacks on deep neural networks by activation clustering," in *AAAI Workshop on Artificial Intelligence Safety*, 2019. [Online]. Available: <https://arxiv.org/abs/1811.03728>
- [34] Y. Gao, C. Xu, D. Wang, S. Chen, D. C. Ranasinghe, and S. Nepal, "STRIP: A defence against trojan attacks on deep neural networks," in *Annual Computer Security Applications Conference (ACSAC)*, 2019.

- [35] B. Tran, J. Li, and A. Madry, "Spectral signatures in backdoor attacks," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2018. [Online]. Available: <https://arxiv.org/abs/1811.00636>
- [36] X. Yang, X. Wang, Q. Zhang, L. Petzold, W. Y. Wang, X. Cheng, and L. Hou, "Shadow alignment: The ease of subverting safely-aligned language models," *arXiv preprint*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.02949>
- [37] P. F. Christiano, J. Leike, T. B. Brown, M. Martic, S. Legg, and D. Amodei, "Deep reinforcement learning from human preferences," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. [Online]. Available: <https://arxiv.org/abs/1706.03741>
- [38] L. Ouyang, J. Wu, X. Jiang *et al.*, "Training language models to follow instructions with human feedback," *Advances in Neural Information Processing Systems*, vol. 35, 2022. [Online]. Available: <https://arxiv.org/abs/2203.02155>
- [39] X. Xu, Q. Wang, H. Li, N. Borisov, C. A. Gunter, and B. Li, "Detecting AI trojans using meta neural analysis," in *IEEE Symposium on Security and Privacy (S&P)*, 2021.
- [40] R. Fu, Z. Zhao, X. Yang, X. Zhang *et al.*, "ShrinkPad: A stealthy and efficient textual backdoor attack framework with pad token," in *Proceedings of the 30th ACM International Conference on Multimedia (MM)*, 2022.
- [41] R. Schaeffer, B. Miranda, and S. Koyejo, "Are emergent abilities of large language models a mirage?" in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. [Online]. Available: <https://arxiv.org/abs/2304.15004>
- [42] Z. Wang *et al.*, "WaDec: Detecting LLM backdoors via weight decomposition," in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024. [Online]. Available: <https://arxiv.org/abs/2407.01011>
- [43] J. Guo *et al.*, "SCALE-UP: An efficient black-box input-level backdoor detection via analyzing scaled prediction consistency," in *International Conference on Learning Representations (ICLR)*, 2023. [Online]. Available: <https://arxiv.org/abs/2302.03251>
- [44] T. A. Nguyen and A. Tran, "WaNet - imperceptible warping-based backdoor attack," in *International Conference on Learning Representations (ICLR)*, 2021. [Online]. Available: <https://arxiv.org/abs/2102.10369>